


JevBOT Rx: architecture

A hospital delivery robot. Code guarantees safety, and Jev supplies fast, bounded judgment. Built at JEVATHON SF on 2026-09-26, on top of JevPilot by Standard Agents.



Where the time goes

Jev returns a probability over the options it's offered, with no text generation, so answers arrive in **≈130 ms**. The LLM writes thinking tokens and then JSON, which takes **≈2.5–5 s**. By the time its answer lands, people and carts have moved, so the robot either waits or acts on a stale view of the hallway.

Where the cost goes

Jev charges for input only (\$0.042/M) and output is free. The LLM pays for input, output, and hidden thinking tokens. Measured per decision: **\$0.00006** vs **\$0.0002**.

Why it's safe

Code does collision prediction, filtering, validation and the safety brake, identically for both brains. The model only picks among paths that were already checked. With the brake on, both robots finished with 0 contacts, and freshness shows up as delivery time.

Measured, seed 42	Jev	DeepSeek-V4-Flash (GMI)
Standard run: Pharmacy → ICU bed 4	84 s	115 s (30 s later)
Time per decision (p50)	≈126 ms	≈2,500 ms
Cost per decision	\$0.00006	\$0.0002
Stress test (safety brake off for both, patient steps into the lane about 2 s ahead)	Mixed: in one run only the LLM hit the patient; in another, both did. Too few runs to claim a rate	

These are single runs on one seed, reported as measured, not averaged. Reproduce them with `race.html?seed=42` and `race.html?seed=42&stress=1`.